CPA and CCA security

CPA & the IND-CPA game

You have an oracle, an adversary who is allowed to keep sending messages to the oracle at any point in time: $m_1, m_2 \ldots m_N$ and the oracle will send back $E(m_1), E(m_2), \ldots, E(m_N)$.

The adversary will send m_1', m_2' (they could be one of the $m_1, m_2 \ldots m_N$) and the oracle encrypts one of them (picks b randomly, 0 or 1) and sends the encryption to the adversary, who's supposed to tell whether m_1' or m_2' was encrypted. Therefore, no CPA scheme can be deterministic, since the adversary could already have the encryption of m_1' or m_2'.

$$\Pr[b' = b] \leq \frac{1}{2} + \varepsilon$$

ε is the advantage here

CCA & the IND-CCA game

The adversary also has access to decryption. He would always win since he would query the decryption oracle with the $E(m_b')$ challenge he received from the encryption oracle. Therefore, the attacker cannot query the decryption oracle with the challenges that he receives from the encryption oracle.

This is the **strongest model** of security for a cryptosystem.

Example of CPA secure system that's not CCA secure

$$Enc_k(m) = \langle r, F_k(r) \oplus m \rangle, where \ |m| = n, r \leftarrow U_n \ and \ F_k \ is \ a \ PRF$$

A CCA attack

Adversary sends $m_0 = 0^n$ and $m_1 = 1^n$ to the encryption oracle.

He gets back $Enc_k(m_b) = \langle r, F_k(r) \oplus m_b \rangle, b \leftarrow u$

- $Enc_k(m_b) = \begin{cases} Enc_k(m_0) = \langle r_0, F_k(r_0) \rangle, 50\% \ of \ the \ time \\ Enc_k(m_1) = \langle r_1, \overline{F_k(r_1)} \rangle, 50\% \ of \ the \ time \end{cases}$

Since the decryption oracle will not accept the encryption of the challenge messages (m_0 or m_1) as input, the attacker will do the next best thing. He can flip the first bit in the encryption of m_b.

The decryption oracle will gladly decrypt the new ciphertext $\langle r, \ c \oplus 10^{n-1} \rangle$.

The adversary can now tell from what he gets back (either 01^{n-1} or 10^{n-1}) whether $m_0 = 0^n$ or $m_1 = 1^n$ was encrypted.

Constructing CCA-secure encryption schemes

Common approach: CPA-ENC + EU-MAC

Most (or maybe all) CCA attacks rely on the attacker modifying $E_k(m'_b)$ and then feeding it to the decryption oracle whose answer will enable the attacker to tell which message was encrypted. If you add a MAC on top of the encryption, disallowing the attacker to fiddle with the ciphertext, then you essentially render the decryption oracle useless.

Construction: Let $\pi = (Enc, Dec, Gen)$. The generator now has to generate two keys, one for the MAC and one for the encryption scheme.

$$Gen(msg\ size) = \langle key_{enc}, key_{mac} \rangle$$

$$Enc(key_{enc}, key_{mac}, msg) = \langle ctxt, mac_{tag} \rangle$$

$$Dec(key_{enc}, key_{mac}, ctxt, tag) = plaintext\ or\ \bot$$

Theorem: If π_E is IND-CPA-secure and π_M is EU-MAC-secure then construction #1 is CCA-secure.

Proof (idea): Queries to the decryption oracle must be of the form $\langle c, t \rangle$, where:

$$c \leftarrow E_{k_1}(m)\ and\ t \leftarrow Mac_{k_2}(c)$$

Since the MAC is secure the attacker has virtually no chance of modifying the challenge ciphertext without the decryption oracle knowing, effectively rendering the oracle useless.

Three ways of implementing EU-MAC and CCA-ENC

Encrypt and authenticate independently
This method computes the MAC of the message, not of the encryption. The ciphertext looks like this:

$$\langle c, t \rangle, where\ c \leftarrow E_{k_{enc}}(m)\ and\ t \leftarrow Mac_{key_{mac}}(m)$$

Counterexample:

π_e is IND-CPA, $c \leftarrow E_{k_1}(m)$

π_M is EU-MAC, $t \leftarrow \left(m, Mac_{k_2}(m) \right)$

This is bad, because the tag t exposes the message m.

$$\langle c, t \rangle = \langle E_{k_{enc}}(m), \langle m, Mac_{key_{mac}}(m) \rangle \rangle$$

Authenticate then encrypt
This method also computes the MAC of the message, but this time the MAC is concatenated with the message m and together they are encrypted to form the ciphertext.

$$\langle c, t \rangle, where\ t \leftarrow Mac_{k_2}(m)\ and\ c \leftarrow E_{k_1}(m \parallel t)$$

Counterexample:

Let $Tr(m) = $ replace each bit b of m with $F(b)$.

Alin Tomescu, CSE408
Tuesday, March 3[rd], Lecture #10

$F(b)$ works very simply. If $b = 0$, then $F(b) = 00$ and if $b = 1$, then $F(b) = 01$ or 10 with equal probability.

$$F(b) = \begin{cases} F(0) = 00, with\ p = 1.0 \\ F(1) = 01, with\ p = 0.5 \\ F(1) = 10, with\ p = 0.5 \end{cases}$$

Let Enc_k be a CTR-mode block-cipher. Let our encryption algorithm be $Enc_k(m) = CTR\big(Tr(m \parallel t)\big)$.

Then, our decryption algorithm will be $Dec_k(\langle c, t \rangle) = Tr^{-1}\big(CTR^{-1}(c)\big)$

If the first bit of the message m is 1 then, after encryption, $\langle c, t \rangle = \begin{cases} CTR(01 \dots) \\ CTR(10 \dots) \end{cases}$

If the first bit of the message m is 0 then, after encryption, $\langle c, t \rangle = CTR(00 \dots)$

Challenge starts, attacker sends $m_0 = 0^n$ and $m_1 = 1^n$ to the encryption oracle which sends the encryption $\langle c, t \rangle = \langle Enc_k(m_b), t \rangle, where\ b \leftarrow u$ of either m_0 or m_1.

The attacker will use the following strategy:

He will **flip the first two bits** in the ciphertext $\langle c, t \rangle$, and send the modified query to the decryption oracle. The decryption oracle will first run CTR mode in reverse, obtaining $Tr(m \parallel t)$. It will then run $Tr^{-1}\big(Tr(m \parallel t)\big)$. There are two cases now. (Let $m = m_b$).

1. The first bit of the message m was 1, then flipping the first two bits in the ciphertext will either change 01 to 10, or 10 to 01 in the transformation $Tr(m \parallel t)$ of the original message m (remember XOR and its properties and the fact that $Enc_k(m) = CTR\big(Tr(m \parallel t)\big)$).
    a. Either way, while decrypting, $Tr^{-1}\big(Tr(m \parallel t)\big)$ will run successfully returning $m \parallel t$.
    b. The decryption oracle will successfully check the integrity of m using the tag t, since there was no overall change in the original message, even though we flipped bits, and it will send m to the attacker.
    c. The attacker now knows m, so he knows which message was encrypted, which breaks the CCA security.
2. The first bit of the message m was 0, then flipping the first two bits in the ciphertext will change 00 to 11 in the transformation $Tr(m \parallel t)$ of the original message m.
    a. This is bad, since there is no reverse mapping from 11 in $F(b)$, which means Tr^{-1} will give an error.
        i. Remember, how $F(b)$ works, it only maps bits to 00, 01 and 10. It doesn't map anything to 11.
    b. Therefore, when this happens the decryption oracle will return $\perp$, indicating it couldn't decrypt the ciphertext.
    c. The attacker now knows that the first bit of the message m was 0, since he knows the only error that was introduced was the flipping of the first two bits in the ciphertext which resulted in the first bits of $Tr(m \parallel t)$ to be set to 11, making Tr^{-1} fail. This breaks the CCA security since the attacker will know m_0 was the encrypted message, since m_1 was all one's.

Encrypt then authenticate
Our first example proved this construction was CCA-secure.

Alin Tomescu, CSE408
Tuesday, March 3rd, Lecture #10

Public-key cryptography

Key distribution is a problem. Key revocation is a problem, if you want to change someone's key, you do that by changing everyone else's key to a new key. Public-key cryptography is here to solve these problems.